

1. Introduction

We present a novel framework for cross-lingual document retrieval evaluated on standard benchmarks. Our approach achieves state-of-the-art results while using fewer parameters than prior work. We further provide a detailed analysis of failure modes and discuss implications for future research. The proposed method transfers to low-resource settings without architectural changes, and we report ablations in the appendix. The results suggest that parameter sharing across languages serves as a strong inductive bias for retrieval tasks. Code and data will be released upon publication.