

Results

Our model achieves a 4.2 point improvement in mean reciprocal rank over the strongest baseline while using approximately 30% fewer parameters and converging in half the training time. Our model achieves a 4.2 point improvement in mean reciprocal rank over the strongest baseline while using approximately 30% fewer parameters and converging in half the training time. Our model achieves a 4.2 point improvement in mean reciprocal rank over the strongest baseline while using approximately 30% fewer parameters and converging in half the training time.